HPLT Analytics report

HPLTAnalytics

General overview

Corpus	Analytics date	Language
HPLT-docslite.lv.tsv	6/8/2024	Latvian (lv)

Volumes

Docs	Segments	Unique segments	Tokens	Size	Characters
1,537,254	185,494,115	43,993,827 (23.72 %)	2B	11.55 GB	

Top 10 domains

Domain	Docs	% of total
diebuchsuche.com	112K	7.29
viss.lv	107K	6.96
delfi.lv	40K	2.58
europages.lv	31K	2.02
tvnet.lv	17K	1.09
lsm.lv	17K	1.08
agoda.com	12K	0.78
maminuklubs.lv	12K	0.75
wikipedia.org	11K	0.72
la.lv	11K	0.71

Top 10 TLDs

Domain	Docs	% of total
lv	1M	67.83
com	315K	20.48
eu	31K	2.03
org	26K	1.67
net	17K	1.13
gov.lv	16K	1.03
lt	14K	0.90
info	12K	0.79
io	6K	0.39
co.uk	4.8K	0.31

Documents size (in segments)

<= 25 segments **8.69%** (133K documents)
> 25 segments **91.31%** (1.4M documents)

Documents by collection

cc40 (530K)
wide16 (524K)
wide17 (179K)
wide15 (303K)

Language Distribution

Number of segments

- Latvian (lv) - 113M
- English (en) - 28M
- French (fr) - 6.4M
- German (de) - 4.6M
- Lithuanian (lt) - 2.8M
- Spanish (es) - 2.7M
- Dutch (nl) - 2.6M
- hbs-lat - 2.3M
- Italian (it) - 2.2M
- Finnish (fi) - 1.7M
- 164 Others - 19M

Percentage of segments in Latvian (lv) inside documents

Distribution of documents by document score

score < 5 - **26.26%** (403K documents)
score >= 5 - **73.74%** (1.1M documents)

Segment length distribution by token

<= 49 tokens = **38M** segments | **140M** duplicates
> 50 tokens = **7.9M** segments | **1.9M** duplicates

Number of tokens in the segment

- Unique segments
- Duplicated segments

Segment noise distribution

Too long	0.25 %
Too short	43.93 %
URLs	3.18 %
Bad encoding	0.00 %

About HPLT Analytics

Volumes - Segments

Segments correspond to paragraph and list boundaries as defined by HTML elements (<p>, <ul>, <ol>, etc.) replaced by newlines.

Volumes - Tokens

Tokenized with https://github.com/hplt-project/data-analytics-tool/blob/main/tokenizers-info.md

Type-Token Ratio

Lexical variety computed as *number or types (uniques)/number of tokens*, after removing punctuation (https://www.sltinfo.com/wp-content/uploads/2014/01/type-token-ratio.pdf).

Document size (in segments)

Segments correspond to paragraph and list boundaries as defined by HTML elements (<p>, <ul>, <ol>, etc.) replaced by newlines.

Language distribution

Language identified with FastSpell (https://github.com/mbanon/fastspell).

Distribution of segments by fluency score

Obtained with Monocleaner (https://github.com/bitextor/monocleaner).

Distribution of documents by average fluency score

Obtained with Monocleaner (https://github.com/bitextor/monocleaner).

Distribution of documents by document score

Obtained with Web Docs Scorer (https://github.com/pablop16n/web-docs-scorer/).

Segment length distribution by token

Tokenized with https://github.com/hplt-project/data-analytics-tool/blob/main/tokenizers-info.md

Segment noise distribution

Obtained with Bicleaner Hardrules (https://github.com/bitextor/bicleaner-hardrules/).

Frequent n-grams

Tokenized with https://github.com/hplt-project/data-analytics-tool/blob/main/tokenizers-info.md, after removing n-grams starting or ending in a stopword. Stopwords from https://github.com/hplt-project/data-analytics-tool/blob/main/scripts/resources/README.txt