HPLT Analytics report

HPLTAnalytics

General overview

Corpus	Analytics date	Language
dik_Latn.jsonl.tsv	10/31/2024	Dinka (dik)

Volumes

Docs	Segments	Unique segments	Tokens	Size	Characters
2,325	34,647	15,747 (45.45 %)	2.9M	12.52 MB	11,505,568

Top 10 domains

Domain	Docs	% of total
bible.is	861	37.03
wikipedia.org	435	18.71
stepbible.org	245	10.54
pngscriptures.org	240	10.32
ebible.org	97	4.17
png.bible	59	2.54
communitydoor.org.au	36	1.55
consumer.vic.gov.au	35	1.51
raizetechnical.com	22	0.95
1800respect.org.au	18	0.77

Top 10 TLDs

Domain	Docs	% of total
org	1.1K	45.94
is	861	37.03
com	88	3.78
org.au	74	3.18
vic.gov.au	72	3.10
bible	59	2.54
gov.au	33	1.42
com.au	15	0.65
in	12	0.52
net	9	0.39

Documents size (in segments)

<= 25 segments **81.68%** (1.9K documents)
> 25 segments **18.32%** (424 documents)

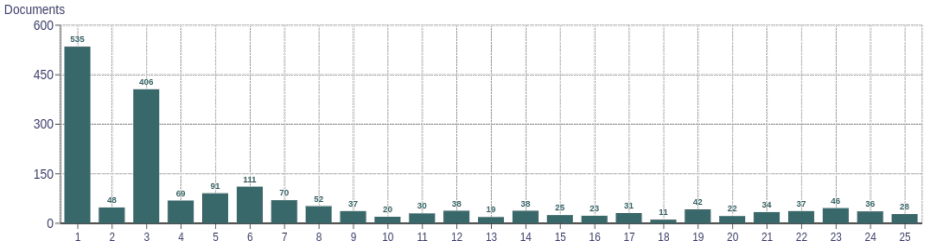

Documents by collection

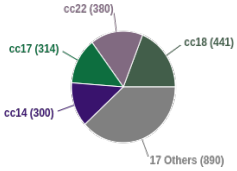

Language Distribution

Number of segments

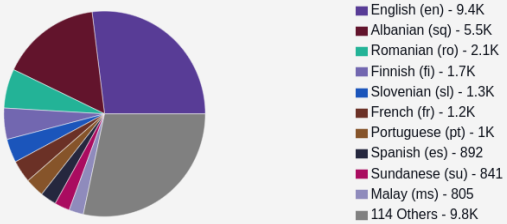

*Dinka (dik) identification might be inaccurate because language is not supported by Fasttext

Percentage of segments in Dinka (dik) inside documents

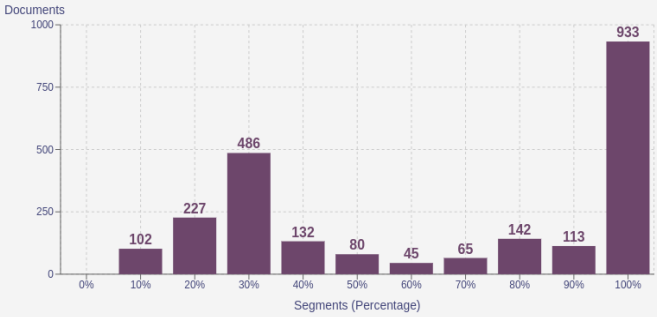

Distribution of documents by document score

score < 5 - **0%** (0 documents)
score >= 5 - **100%** (2.3K documents)

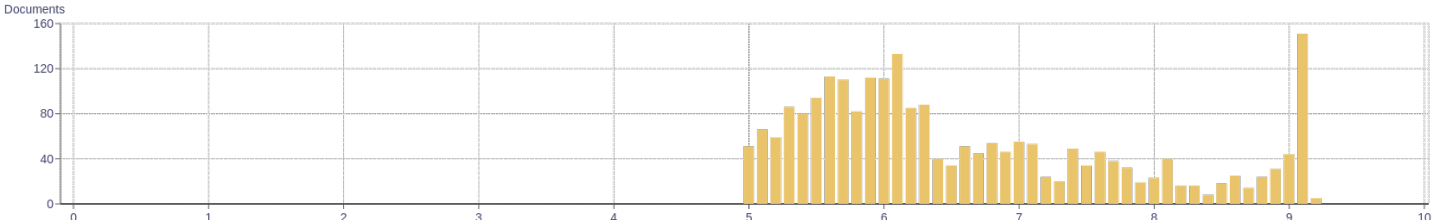

Segment length distribution by token

<= 49 tokens = **11K** segments | **15K** duplicates
> 50 tokens = **8.8K** segments | **3.8K** duplicates

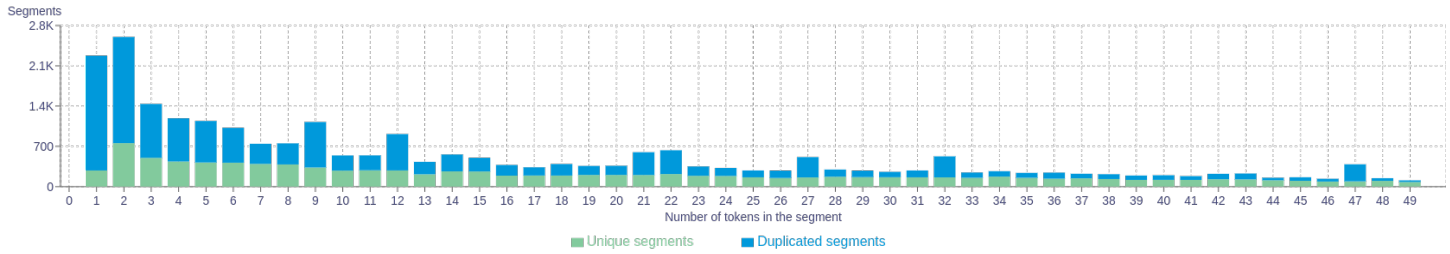

Segment noise distribution

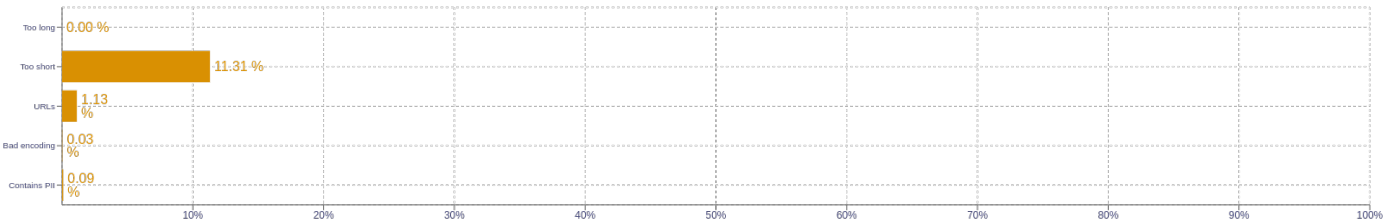

About HPLT Analytics

Volumes - Segments

Segments correspond to paragraph and list boundaries as defined by HTML elements (<p>, <ul>, <ol>, etc.) replaced by newlines.

Volumes - Tokens

Tokenized with https://github.com/hplt-project/data-analytics-tool/blob/main/tokenizers-info.md

Type-Token Ratio

Lexical variety computed as *number or types (uniques)/number of tokens*, after removing punctuation (https://www.sltinfo.com/wp-content/uploads/2014/01/type-token-ratio.pdf).

Document size (in segments)

Segments correspond to paragraph and list boundaries as defined by HTML elements (<p>, <ul>, <ol>, etc.) replaced by newlines.

Language distribution

Language identified with FastSpell (https://github.com/mbanon/fastspell).

Distribution of segments by fluency score

Obtained with Monocleaner (https://github.com/bitextor/monocleaner).

Distribution of documents by average fluency score

Obtained with Monocleaner (https://github.com/bitextor/monocleaner).

Distribution of documents by document score

Obtained with Web Docs Scorer (https://github.com/pablop16n/web-docs-scorer/).

Segment length distribution by token

Tokenized with https://github.com/hplt-project/data-analytics-tool/blob/main/tokenizers-info.md

Segment noise distribution

Obtained with Bicleaner Hardrules (https://github.com/bitextor/bicleaner-hardrules/).

Frequent n-grams

Tokenized with https://github.com/hplt-project/data-analytics-tool/blob/main/tokenizers-info.md, after removing n-grams starting or ending in a stopword. Stopwords from https://github.com/hplt-project/data-analytics-tool/blob/main/scripts/resources/README.txt